

In [1]:

```
#Часть 1
import numpy as np
import pandas as pd
```

In [2]:

```
data = pd.read_csv('adult.data.csv')
data.head()
```

Out[2]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

In [3]:

```
#1. How many men and women (sex feature) are represented in this dataset?
data["sex"].value_counts()
```

Out[3]:

Male 21790
Female 10771
Name: sex, dtype: int64

In [4]:

```
#2. What is the average age (age feature) of women?
data.loc[data["sex"] == 'Female', 'age'].mean()
```

Out[4]:

36.85823043357163

In [5]:

```
#3. What is the proportion of German citizens (native-country feature)?
print("{0:%}".format(data[data["native-country"] == "Germany"].shape[0] / data.shape[0]))
```

0.420749%

In [6]:

```
#4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (salary feature)
#and those who earn less than 50K per year?
```

```
ages1 = data[data["salary"] == "<=50K"]["age"]
ages2 = data[data["salary"] == ">50K"]["age"]
print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))
print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))
```

< 50K: 36.782372786407767 ± 14.020088400824813 years

$\leq 50K$: = 36.76373766407767 \pm 14.020068490824613 years
 $> 50K$: = 44.24984058155847 \pm 10.51902771985177 years

In [7]:

```
#6. Is it true that people who earn more than 50K have at least high school education?
#(education – Bachelors, Prof-school, Assoc-acdm, Assoc-voc, Masters or Doctorate feature)
high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm",
                       "Assoc-voc", "Masters", "Doctorate"])
def high_educated(e):
    return e in high_educations

data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

Out[7]:

False

In [8]:

```
#7. Display statistics of age for each race (race feature) and each gender. Use groupby() and describe().
#Find the maximum age of men of Amer-Indian-Eskimo race.
for (race, sex), sub_df in data.groupby(['race', 'sex']):
    print("Race: {0}, sex: {1}".format(race, sex))
    print(sub_df['age'].describe())
```

Race: Amer-Indian-Eskimo, sex: Female

count 119.000000
mean 37.117647
std 13.114991
min 17.000000
25% 27.000000
50% 36.000000
75% 46.000000
max 80.000000

Name: age, dtype: float64

Race: Amer-Indian-Eskimo, sex: Male

count 192.000000
mean 37.208333
std 12.049563
min 17.000000
25% 28.000000
50% 35.000000
75% 45.000000
max 82.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Female

count 346.000000
mean 35.089595
std 12.300845
min 17.000000
25% 25.000000
50% 33.000000
75% 43.750000
max 75.000000

Name: age, dtype: float64

Race: Asian-Pac-Islander, sex: Male

count 693.000000
mean 39.073593
std 12.883944
min 18.000000
25% 29.000000
50% 37.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Black, sex: Female

count 1555.000000
mean 37.854019
std 12.637197
min 17.000000
25% 28.000000
50% 37.000000
75% 46.000000
max 90.000000

Name: age, dtype: float64

Race: Black, sex: Male

count 1569.000000
mean 37.682600
std 12.882612

```

min      17.000000
25%      27.000000
50%      36.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: Other, sex: Female
count    109.000000
mean     31.678899
std      11.631599
min      17.000000
25%      23.000000
50%      29.000000
75%      39.000000
max      74.000000
Name: age, dtype: float64
Race: Other, sex: Male
count    162.000000
mean     34.654321
std      11.355531
min      17.000000
25%      26.000000
50%      32.000000
75%      42.000000
max      77.000000
Name: age, dtype: float64
Race: White, sex: Female
count   8642.000000
mean     36.811618
std      14.329093
min      17.000000
25%      25.000000
50%      35.000000
75%      46.000000
max      90.000000
Name: age, dtype: float64
Race: White, sex: Male
count   19174.000000
mean     39.652498
std      13.436029
min      17.000000
25%      29.000000
50%      38.000000
75%      49.000000
max      90.000000
Name: age, dtype: float64

```

In [9]:

```

#8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (marital-status feature)?
#Consider as married those who have a marital-status starting with Married (Married-civ-spouse, Married-spouse-absent or
#Married-AF-spouse), the rest are considered bachelors.

```

```

data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].isin(['Never-married',
                                       'Separated',
                                       'Divorced',
                                       'Widowed']))], 'salary'].value_counts()

```

Out[10]:

```

<=50K    7552

```

File "<ipython-input-9-12be563d53cd>", line 10

Out[10]:

^

SyntaxError: invalid syntax

In []:

```

data.loc[(data['sex'] == 'Male') &
         (data['marital-status'].str.startswith('Married'))], 'salary'].value_counts()

```

In []:

```

data['marital-status'].value_counts()

```

In []:

```
#9. What is the maximum number of hours a person works per week (hours-per-week feature)?
#How many people work such a number of hours and what is the percentage of those who earn a lot among them?
max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))

num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))

rich_share = float(data[(data['hours-per-week'] == max_load)
    & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_share)))
```

In []:

```
#10. Count the average time of work (hours-per-week) those who earning a little and a lot (salary) for each country
#(native-country).
pd.crosstab(data['native-country'], data['salary'],
    values=data['hours-per-week'], aggfunc=np.mean).T
```

In [10]:

```
#Часть 2
!pip install pandasql
```

Requirement already satisfied: pandasql in c:\users\dovlat\anaconda3\lib\site-packages (0.7.3)
Requirement already satisfied: sqlalchemy in c:\users\dovlat\anaconda3\lib\site-packages (from pandasql) (1.3.1)
Requirement already satisfied: pandas in c:\users\dovlat\anaconda3\lib\site-packages (from pandasql) (0.24.2)
Requirement already satisfied: numpy in c:\users\dovlat\anaconda3\lib\site-packages (from pandasql) (1.16.2)
Requirement already satisfied: python-dateutil>=2.5.0 in c:\users\dovlat\anaconda3\lib\site-packages (from pandas->pandasql) (2.8.0)
Requirement already satisfied: pytz>=2011k in c:\users\dovlat\anaconda3\lib\site-packages (from pandas->pandasql) (2018.9)
Requirement already satisfied: six>=1.5 in c:\users\dovlat\anaconda3\lib\site-packages (from python-dateutil>=2.5.0->pandas->pandasql) (1.12.0)

In [11]:

```
from pandasql import sqldf
pysqldf = lambda q: sqldf(q, globals())
```

In [12]:

```
user_usage = pd.read_csv('user_usage.csv')
```

In [13]:

```
user_device = pd.read_csv('user_device.csv')
```

In [14]:

```
user_usage.head()
```

Out[14]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id
0	21.97	4.82	1557.33	22787
1	1710.08	136.88	7267.55	22788
2	1710.08	136.88	7267.55	22789
3	94.46	35.17	519.12	22790
4	71.59	79.26	1557.33	22792

In [15]:

```
user_usage.dtypes
```

Out[15]:

```
outgoing_mins_per_month    float64
outgoing_sms_per_month     float64
monthly_mb                 float64
use_id                    int64
dtype: object
```

In [16]:

```
user_device.head()
```

Out[16]:

	use_id	user_id	platform	platform_version	device	use_type_id
0	22782	26980	ios	10.2	iPhone7,2	2
1	22783	29628	android	6.0	Nexus 5	3
2	22784	28473	android	5.1	SM-G903F	1
3	22785	15200	ios	10.2	iPhone7,2	3
4	22786	28239	android	6.0	ONE E1003	1

In [17]:

```
user_device.dtypes
```

Out[17]:

use_id int64
user_id int64
platform object
platform_version float64
device object
use_type_id int64
dtype: object

In [18]:

```
user_device.merge(user_usage).head()
```

Out[18]:

	use_id	user_id	platform	platform_version	device	use_type_id	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb
0	22787	12921	android	4.3	GT-I9505	1	21.97	4.82	1557.33
1	22788	28714	android	6.0	SM-G930F	1	1710.08	136.88	7267.55
2	22789	28714	android	6.0	SM-G930F	1	1710.08	136.88	7267.55
3	22790	29592	android	5.1	D2303	1	94.46	35.17	519.12
4	22792	28217	android	5.1	SM-G361F	1	71.59	79.26	1557.33

In [19]:

```
%%timeit  
user_device.merge(user_usage).head()
```

15.7 ms ± 978 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

In [20]:

```
pysqlIdf("""SELECT u.outgoing_mins_per_month, u.outgoing_sms_per_month, u.monthly_mb, u.use_id, d.user_id,d.platform,  
d.platform_version, d.device, d.use_type_id  
FROM user_usage AS u JOIN user_device AS d  
ON u.use_id = d.use_id  
""").head()
```

Out[20]:

	outgoing_mins_per_month	outgoing_sms_per_month	monthly_mb	use_id	user_id	platform	platform_version	device	use_type_id
0	21.97	4.82	1557.33	22787	12921	android	4.3	GT-I9505	1
1	1710.08	136.88	7267.55	22788	28714	android	6.0	SM-G930F	1
2	1710.08	136.88	7267.55	22789	28714	android	6.0	SM-G930F	1
3	94.46	35.17	519.12	22790	29592	android	5.1	D2303	1

In [21]:

```
%%timeit
pysqldb("""SELECT u.outgoing_mins_per_month, u.outgoing_sms_per_month, u.monthly_mb, u.use_id, d.user_id,d.platform,
d.platform_version, d.device, d.use_type_id
FROM user_usage AS u JOIN user_device AS d
ON u.use_id = d.use_id
""")
```

59.6 ms ± 8.11 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

In [22]:

```
user_usage.groupby("use_id")["monthly_mb"].mean().head()
```

Out[22]:

```
use_id
22787  1557.33
22788  7267.55
22789  7267.55
22790   519.12
22792  1557.33
Name: monthly_mb, dtype: float64
```

In [23]:

```
%%timeit
user_usage.groupby("use_id")["monthly_mb"].mean()
```

2.84 ms ± 273 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)